

Under What Conditions Does Cooperation Emerge Among Self-Interested Agents?

An Illustrative Agent-Based Study with the Prisoner’s Dilemma

Scott Brodie Forsyth

July 3, 2026

Abstract

Self-interested agents—individuals who rationally pursue their own payoffs—face a puzzle: in many one-shot social dilemmas, mutual defection is the unique Nash equilibrium, yet cooperation is widespread in nature and human societies. This paper synthesizes canonical game-theoretic and evolutionary mechanisms and supplements them with reproducible Python simulations. The analysis shows, both analytically and computationally, that cooperation can *emerge and persist* only when the strategic environment satisfies additional structure: repeated interaction with a sufficiently high discount factor, spatial or network localization, positive assortment (including kinship), or institutions that alter incentives. The simulations are **illustrative**: they demonstrate comparative statics consistent with established theory but do **not** identify causal effects in observational data, nor do they substitute for laboratory or field experiments. The contribution is pedagogical and methodological—a transparent bridge between formal models and visual intuition for readers without specialist training.

1 Introduction

Imagine two farmers sharing an irrigation canal. Each can contribute labor to maintain it or free-ride on the other’s effort. If they interact once and never meet again, standard game theory predicts both will shirk: defection is a dominant strategy in the one-shot Prisoner’s Dilemma (PD). Yet real communities often cooperate—sometimes reliably, sometimes fragily. The central question is therefore not *whether* self-interest exists, but **under what conditions self-interest leads to cooperation rather than mutual harm**.

This study addresses that question at two levels. First, it states precise mathematical conditions from repeated-game theory and evolutionary biology. Second, it implements agent-based models in Python to visualize how cooperation rises or collapses when those conditions change. Throughout, three types of claims are kept distinct:

1. **Theoretical claims** (e.g., folk-theorem logic): valid given stated assumptions.
2. **Simulation patterns**: numerical illustrations of theory, not independent empirical proof.
3. **Empirical generalizations** about humans: supported by literatures cited here but not replicated in this study.

2 The One-Shot Prisoner’s Dilemma

2.1 Payoffs and the social dilemma

Two players simultaneously choose Cooperate (C) or Defect (D). Payoffs satisfy $T > R > P > S$ and $2R > T + S$, where T is temptation, R reward, P punishment, and S sucker’s payoff. The baseline calibration is $T = 5$, $R = 3$, $P = 1$, $S = 0$ (Figure 1).




Figure 1: Stage-game payoff matrix. Mutual cooperation yields (3, 3); unilateral defection yields (5, 0) or (0, 5); mutual defection yields (1, 1).

The payoff matrix is

$$\Pi = \begin{pmatrix} R & S \\ T & P \end{pmatrix}, \quad (1)$$

where rows index the focal player's action and columns the opponent's.

2.2 Nash equilibrium in one shot

Defection strictly dominates cooperation: regardless of the opponent's action, D yields higher payoff ($T > R$ and $P > S$). Hence the unique pure-strategy Nash equilibrium is (D, D) with payoffs (P, P) , which is Pareto-dominated by (C, C) . This is the canonical *social dilemma*: individual rationality conflicts with collective welfare (Hardin, 1968; Binmore, 1994).

Plain-language summary. If people meet once, never interact again, and care only about immediate payoff, cooperation should unravel. Real life often violates at least one of these premises—and that is where cooperation becomes possible.

3 Mechanisms That Enable Cooperation

Following Nowak (2006) and West et al. (2007), five mechanism families organize the discussion. Table 1 maps each to a qualifying condition.

Table 1: Five routes to cooperation among self-interested agents (Nowak, 2006)

Mechanism	Qualifying condition	Intuition
Kin selection	Relatedness r high; $rb > c$	Help copies of own genes
Direct reciprocity	Repeated play; high δ	Today's help buys tomorrow's help
Indirect reciprocity	Reputation observability	Cooperators gain social standing
Network reciprocity	Local interaction	Cooperators cluster and protect each other
Group selection	Between-group variance	Groups of cooperators outcompete defectors

3.1 Direct reciprocity and repeated games

When the PD is repeated with discount factor $\delta \in (0, 1)$, players can condition future play on past actions. Consider *grim trigger*: cooperate until the opponent defects, then defect forever.

Mutual cooperation is a subgame-perfect equilibrium if the stream from cooperation exceeds the one-shot gain from defection:

$$\frac{R}{1-\delta} \geq T + \frac{\delta P}{1-\delta}. \quad (2)$$

Rearranging yields the sharp threshold

$$\delta \geq \delta^* \equiv \frac{T-R}{T-P}. \quad (3)$$

Proposition 1 (Grim-trigger cooperation threshold). *Under perfect monitoring and the payoffs in (1), mutual cooperation sustained by grim trigger is a subgame-perfect equilibrium of the infinitely repeated PD if and only if $\delta \geq (T-R)/(T-P)$.*

Proof sketch. If both players cooperate, each receives $R/(1-\delta)$. A unilateral deviation yields T in the current period and P in every subsequent period, totaling $T + \delta P/(1-\delta)$. Cooperation is optimal when $R/(1-\delta) \geq T + \delta P/(1-\delta)$, which rearranges to (3). The converse holds by the same comparison. See [Fudenberg and Tirole \(1991\)](#) for the general folk-theorem framework. $\square$

For the baseline payoffs, $\delta^* = (5-3)/(5-1) = 0.5$. Figure 2 plots this feasibility region.




Figure 2: Grim-trigger cooperation is sustainable for $\delta \geq \delta^*$. Below the threshold, defecting once and returning to mutual punishment is preferable to continued cooperation.

[Axelrod and Hamilton \(1981\)](#) showed in computer tournaments that *Tit-for-Tat*—cooperate first, then mirror the opponent—performs well in repeated PDs with uncertain horizon. However, [Press and Dyson \(2012\)](#) later showed that extortionate strategies can dominate in evolutionary settings, indicating that “nice” reciprocity is not the only equilibrium outcome. This study does not resolve equilibrium selection; it characterizes only *when cooperation is possible*.

3.2 Kin selection and positive assortment

Hamilton’s rule states that a costly cooperative act can spread by natural selection when

$$rb > c, \quad (4)$$

where r is genetic relatedness between actor and recipient, b is benefit to recipient, and c is cost to actor ([Hamilton, 1964](#); [Trivers, 1971](#)). Mapping the PD to a donation game with benefit $b = R - P$ and cost $c = T - R$ yields the threshold $r > c/b = (T-R)/(R-P)$. Under the baseline payoffs, $r^* = 2/2 = 1$: kin selection alone requires maximal relatedness unless payoffs are rescaled. *Positive assortment*—non-random pairing of cooperators—plays an analogous role in behavioral and cultural evolution ([Lehmann and Keller, 2006](#)).

Figure 3 reports a stylized group model in which cooperation rises with an assortment parameter r . The dashed vertical line marks the Hamilton threshold under the donation framing; the simulation parameter is a behavioral assortment rate, not a direct estimate of genetic relatedness, so the line is a theoretical benchmark rather than a predicted breakpoint.




Figure 3: Cooperation fraction vs. assortment in a group-structured simulation (10 replicates per point). The dashed line marks the Hamilton threshold $r^* = c/b = 1$ under the baseline payoff calibration (donation framing).

3.3 Spatial and network reciprocity

When interactions are *local*—on a lattice or network—cooperators can form clusters that earn high payoffs from each other while isolating exploiters (Nowak and May, 1992; Szabó and Fátih, 2007). Nowak and May (1992) showed that spatial structure can sustain cooperation even when well-mixed populations cannot.

Figures 4–5 contrast spatial and well-mixed dynamics.




Figure 4: Snapshots of a spatial evolutionary PD on a 60×60 toroidal lattice. Cooperators (blue) and defectors (red) self-organize into clusters; boundaries are dynamic.

3.4 Institutions and punishment

Self-interest can also support cooperation when rules change payoffs. Ostrom (1990) documented how communities govern commons through monitoring, graduated sanctions, and conflict-




Figure 5: Left: higher strategy noise (ε) erodes cooperation. Right: identical payoffs but different population structure yield opposite long-run outcomes in a representative run (seed fixed).

resolution mechanisms—not through naive altruism. Laboratory evidence shows that *altruistic punishment* of defectors can sustain cooperation in public-goods games (Fehr and Gächter, 2002), though punishment itself is costly and culturally variable (Henrich et al., 2001). Institutional design is not simulated here; it is noted as a distinct, empirically important channel.

4 Computational Methods

All code is in `src/simulation.py` and `src/generate_figures.py`. Models include:

- **Spatial PD:** 50×50 toroidal lattice, von Neumann neighbors, synchronous imitate-best-neighbor update with noise $\varepsilon \in [0, 1]$.
- **Well-mixed Moran process:** fitness-proportional imitation with mutation rate 0.01.
- **Group assortment:** reproduction with probability r of copying a randomly chosen within-group member.
- **Parameter sweeps:** temptation T , noise ε , assortment r ; 10–12 independent replicates per grid point; error bars are approximate 95% confidence intervals ($\pm 1.96 \times \text{SE}$).

Random seeds are fixed for reproducibility. Simulations run for 250–1500 generations depending on the model. Outputs are **not** calibrated to field data; they illustrate theoretical comparative statics.

5 Results: Comparative Statics

Figure 6 varies temptation T while holding $R = 3$, $P = 1$, $S = 0$. As defection becomes more lucrative, cooperation eventually collapses—consistent with the comparative static that higher T shrinks the basin of attraction for cooperative clusters (Nowak and May, 1992).

Figure 7 shows that execution errors and mutation (ε) reduce cooperation on average, consistent with the idea that reciprocity requires sufficiently accurate memory and implementation (Nowak et al., 2010). The relationship need not be strictly monotonic at every point because of stochastic cluster dynamics.

Figure 8 aggregates illustrative outcomes across mechanisms. One-shot play and well-mixed evolution yield low cooperation; repeated-game feasibility (here $\delta = 0.8 > \delta^*$), spatial structure, and positive assortment yield high cooperation *under the parameterizations used here*.




Figure 6: Final cooperation vs. temptation T in the spatial model (12 replicates per point).




Figure 7: Noise erodes cooperative clusters by introducing defector invasions (10 replicates per point).




Figure 8: Illustrative cooperation levels by mechanism. Spatial, well-mixed, and assortment bars come from single-seed simulation runs; the one-shot bar is the Nash prediction (0); the repeated-game bar marks grim-trigger feasibility at $\delta = 0.8$.

6 Scope, Claims, and Limitations

6.1 What this study establishes

1. **Logical necessity in models:** In the one-shot PD, defection is dominant; cooperation requires modifying the game (repetition, assortment, structure, or institutions).
2. **Quantitative thresholds:** Equations (3) and (4) give explicit parameter boundaries for two canonical mechanisms.
3. **Qualitative simulation patterns:** Spatial localization and low noise tend to promote cooperation; high temptation and well-mixed interaction tend to suppress it—patterns aligned with [Nowak and May \(1992\)](#); [Nowak \(2006\)](#).

6.2 Limitations

1. **No empirical identification.** The simulations are *stylized models*, not estimates of causal effects in field data. No claim is made that, for example, “raising relatedness by 0.1 *causes* a 15% cooperation increase” in human populations.
2. **Binary strategies.** Real behavior is continuous (partial contributions, negotiation, hybrid norms). Binary C/D is a deliberate simplification.
3. **Equilibrium selection.** Many strategy profiles can coexist in repeated and spatial games; the imitative dynamics implemented here select particular trajectories without general uniqueness ([Fudenberg and Tirole, 1991](#); [Press and Dyson, 2012](#)).
4. **Parameter dependence.** Spatial cooperation persists for some (T, R, P, S) combinations and vanishes for others ([Nowak and May, 1992](#)). The sweeps cover a subset of parameter space only.
5. **Human motives.** Experiments document fairness, guilt, and parochialism beyond material PD payoffs ([Henrich et al., 2001](#); [Rand and Nowak, 2013](#)). Homo economicus is a modeling device, not a complete psychological account.
6. **External validity.** Results from 50×50 lattices may not transfer to large heterogeneous networks, mobile populations, or institutional settings with courts and contracts.

6.3 Plain-language takeaway

Cooperation among self-interested agents is **not mysterious**, but it is **not automatic**. It emerges when the shadow of the future is long enough, when cooperators can find each other (genetically, spatially, or socially), when errors are small enough to preserve reciprocity, or when communities build rules that reward contribution and penalize free-riding. Where those conditions fail—one-shot anonymous interactions with high temptation to defect—self-interest predicts the tragedy of the commons ([Hardin, 1968](#)).

7 Conclusion

This study maps the conditions under which cooperation can emerge from self-interest: repeated interaction (with $\delta \geq (T - R)/(T - P)$), positive assortment (Hamilton’s $rb > c$), spatial or network structure, and institutional enforcement. Python agent-based models illustrate these comparative statics with transparent, reproducible code. The analysis is intentionally bounded: it clarifies *when* and *why* cooperation is theoretically sustainable, while leaving empirical measurement of those conditions in specific societies to field work and experiments ([Ostrom, 1990](#); [Fehr and Gächter, 2002](#); [Henrich et al., 2001](#)).

Replication

```
pip install -r requirements.txt
python src/validate.py
python src/generate_figures.py
cd paper && pdflatex cooperation_emergence && bibtex cooperation_emergence
pdflatex cooperation_emergence && pdflatex cooperation_emergence
```

References

- Axelrod, R. and Hamilton, W. D. (1981). The evolution of cooperation. *Science*, 211(4489):1390–1396.
- Binmore, K. (1994). *Game Theory and the Social Contract, Volume I: Playing Fair*. MIT Press, Cambridge, MA.
- Fehr, E. and Gächter, S. (2002). Altruistic punishment in humans. *Nature*, 415(6868):137–140.
- Fudenberg, D. and Tirole, J. (1991). *Game Theory*. MIT Press, Cambridge, MA.
- Hamilton, W. D. (1964). The genetical evolution of social behaviour. I. *Journal of Theoretical Biology*, 7(1):1–16.
- Hardin, G. (1968). The tragedy of the commons. *Science*, 162(3859):1243–1248.
- Henrich, J., Boyd, R., Bowles, S., Camerer, C., Fehr, E., Gintis, H., et al. (2001). In search of homo economicus: Behavioral experiments in 15 small-scale societies. *American Economic Review*, 91(2):73–78.
- Lehmann, L. and Keller, L. (2006). The evolution of cooperation and altruism—a general framework and a classification of models. *Journal of Evolutionary Biology*, 19(5):1365–1376.
- Nowak, M. A. (2006). Five rules for the evolution of cooperation. *Science*, 314(5805):1560–1563.
- Nowak, M. A. and May, R. M. (1992). Evolutionary games and spatial chaos. *Nature*, 359(6398):826–829.
- Nowak, M. A., Tarnita, C. E., and Antal, T. (2010). Evolutionary dynamics in structured populations. *Philosophical Transactions of the Royal Society B*, 365(1537):19–30.
- Ostrom, E. (1990). *Governing the Commons: The Evolution of Institutions for Collective Action*. Cambridge University Press, Cambridge.
- Press, W. H. and Dyson, F. J. (2012). Iterated prisoner’s dilemma contains strategies that dominate any evolutionary opponent. *Proceedings of the National Academy of Sciences*, 109(26):10409–10413.
- Rand, D. G. and Nowak, M. A. (2013). Human cooperation. *Trends in Cognitive Sciences*, 17(8):413–425.
- Szabó, G. and Fáth, G. (2007). Evolutionary games on graphs. *Physics Reports*, 446(4-6):97–216.
- Trivers, R. L. (1971). The evolution of reciprocal altruism. *The Quarterly Review of Biology*, 46(1):35–57.
- West, S. A., Griffin, A. S., and Gardner, A. (2007). Evolutionary explanations for cooperation. *Current Biology*, 17(16):R661–R672.